

DSC540-T303-Data-Preparation-Week3-4

December 16, 2024

```
[173]: # Weeks 3 & 4 Exercises
```

```
[174]: # 1. The Data Wrangling Workshop: Activity 3.01, page 155
# Generating statistics from a csv file.
```

```
#Load the necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
[175]: #Read the Boston housing dataset from the csv file.
```

```
[176]: df=pd.read_csv("Boston_housing.csv")
```

```
[177]: #Check the first 10 records
```

```
[178]: df.head(10)
```

```
[178]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	
5	0.02985	0.0	2.18	0	0.458	6.430	58.7	6.0622	3	222	18.7	
6	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2	
7	0.14455	12.5	7.87	0	0.524	6.172	96.1	5.9505	5	311	15.2	
8	0.21124	12.5	7.87	0	0.524	5.631	100.0	6.0821	5	311	15.2	
9	0.17004	12.5	7.87	0	0.524	6.004	85.9	6.5921	5	311	15.2	

	B	LSTAT	PRICE
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2
5	394.12	5.21	28.7
6	395.60	12.43	22.9

```

7  396.90  19.15  27.1
8  386.63  29.93  16.5
9  386.71  17.10  18.9

```

```
[179]: #Find the total number of records
```

```
[180]: df.shape
```

```
[180]: (506, 14)
```

```
[181]: #Create a smaller dataset which exclude CHAS, NOX, B, LSTAT
```

```
[182]: df1=df[['CRIM','ZN','INDUS','RM','AGE','DIS',
             'RAD','TAX','PTRATIO','PRICE']]
```

```
[183]: # Check the last 7 records of the new DataFrame you just created
```


```
[184]: df1.tail(7)
```


```
[184]:
```


	CRIM	ZN	INDUS	RM	AGE	DIS	RAD	TAX	PTRATIO	PRICE
499	0.17783	0.0	9.69	5.569	73.5	2.3999	6	391	19.2	17.5
500	0.22438	0.0	9.69	6.027	79.7	2.4982	6	391	19.2	16.8
501	0.06263	0.0	11.93	6.593	69.1	2.4786	1	273	21.0	22.4
502	0.04527	0.0	11.93	6.120	76.7	2.2875	1	273	21.0	20.6
503	0.06076	0.0	11.93	6.976	91.0	2.1675	1	273	21.0	23.9
504	0.10959	0.0	11.93	6.794	89.3	2.3889	1	273	21.0	22.0
505	0.04741	0.0	11.93	6.030	80.8	2.5050	1	273	21.0	11.9

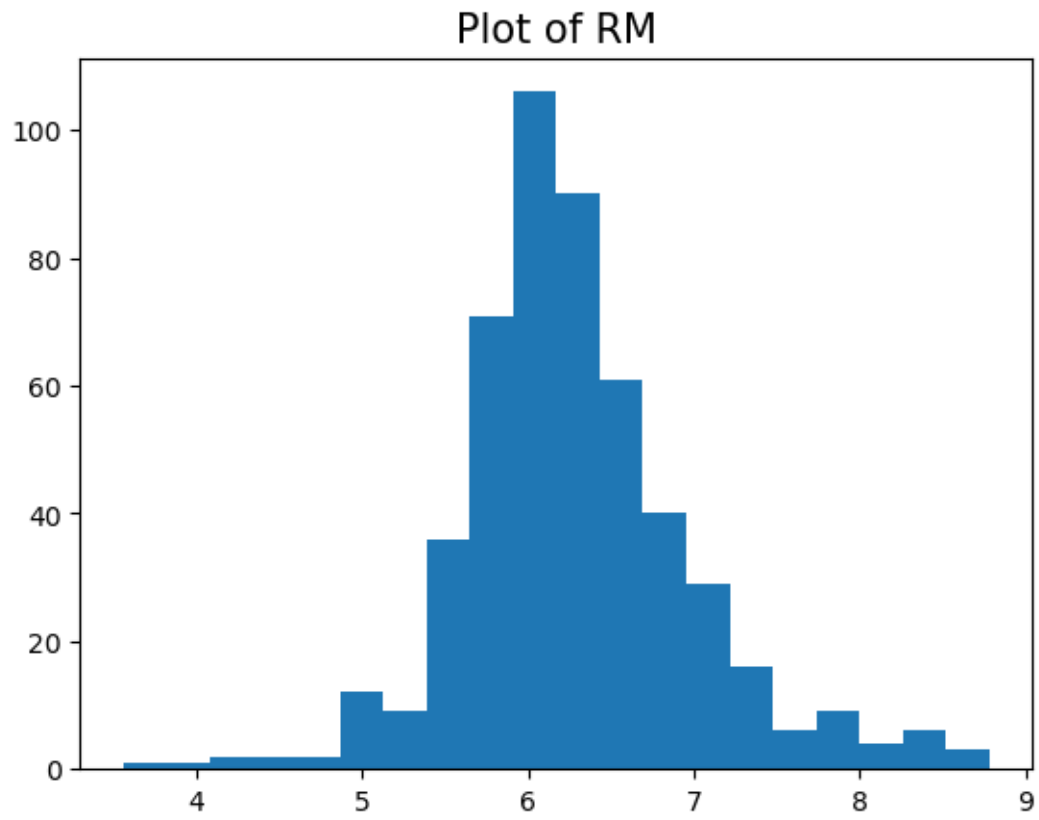
```
[185]: #Plot hs(columns) in the new DataFrame by using a for loop
        histograms of all the
        ↪ variables(columns) in the new DataFrame by using a for loop
```

```
[186]: for c in df1.columns:
        plt.title("Plot of "+c,fontsize=15)
        plt.hist(df1[c],bins=20)
        plt.show()
```













Plot of AGE





Plot of DIS













```
[187]: # Create a scatter plot of crime rate versus price.
```

```
[188]: plt.scatter(df1['CRIM'],df1['PRICE'])  
plt.show()
```



```
[189]: #Create a plot of log10 (crime) veru price.
```

```
[190]: plt.scatter(np.log10(df1['CRIM']),df1['PRICE'],c='red')
plt.title("Crime rate (Log) vs. Price plot", fontsize=18)
plt.xlabel("Log of Crime rate",fontsize=15)
plt.ylabel("Price",fontsize=15)
plt.grid(True)
plt.show()
```



```
[191]: #Calculate the mean room per dwelling
```

```
[192]: df1['RM'].mean()
```

```
[192]: 6.284634387351779
```

```
[193]: #Calculate the median age
```

```
[194]: df1['AGE'].median()
```

```
[194]: 77.5
```

```
[195]: # Calculate the average mean distances to five Bostone employment centers
```

```
[196]: df1['DIS'].mean()
```

```
[196]: 3.795042687747036
```

```
[197]: #Calculate the price of the housing that's than 20
```

```
[198]: low_price=df1['PRICE']<20
print(low_price)
```

```
0      False
1      False
2      False
3      False
4      False
...
501     False
502     False
503     False
504     False
505       True
Name: PRICE, Length: 506, dtype: bool
```

```
[199]: #Calculate the mean of this array
```

```
[200]: # That many houses are priced below 20,000. So that is the answer.
low_price.mean()
```

```
[200]: 0.4150197628458498
```

```
[201]: #Calculate the percentage of houses with a low price < $20000
```

```
[202]: # You can convert that into percentage by multiplying with 100
pcnt=low_price.mean()*100
print("\nPercentage of house with <20,000 price is: ",pcnt)
```

```
Percentage of house with <20,000 price is: 41.50197628458498
```

```
[203]: # 2. The Data Wrangling Workshop: Activity 4.01, page 233
```

```
[204]: # Read the adult income dataset
df = pd.read_csv("adult_income_data.csv")
df.head()
```

```
[204]:   39      State-gov   77516  Bachelors  13      Never-married  \
0  50  Self-emp-not-inc  83311  Bachelors  13  Married-civ-spouse
1  38      Private  215646    HS-grad   9      Divorced
2  53      Private  234721     11th    7  Married-civ-spouse
3  28      Private  338409  Bachelors  13  Married-civ-spouse
4  37      Private  284582    Masters  14  Married-civ-spouse

      Adm-clerical  Not-in-family    Male  2174  0  40  United-States  \
0  Exec-managerial      Husband    Male     0  0  13  United-States
1  Handlers-cleaners  Not-in-family    Male     0  0  40  United-States
```

2	Handlers-cleaners	Husband	Male	0	0	40	United-States
3	Prof-specialty	Wife	Female	0	0	40	Cuba
4	Exec-managerial	Wife	Female	0	0	40	United-States

```

    <=50K
0    <=50K
1    <=50K
2    <=50K
3    <=50K
4    <=50K

```

```
[205]: #Create a script that will read a text file line by line and exaxtract the first
       ↪line which is the header of the .csv file
```

```
[206]: names = []
       with open('adult_income_names.txt','r') as f:
           for line in f:
               f.readline()
               var=line.split(":")[0]
               names.append(var)
       names
```

```
[206]: ['age',
       'workclass',
       'fnlwgt',
       'education',
       'education-num',
       'marital-status',
       'occupation',
       'relationship',
       'sex',
       'capital-gain',
       'capital-loss',
       'hours-per-week',
       'native-country']
```

```
[207]: # Add the name of "Income" for the response variable(last column) to the
       ↪dataset by using the "append" comman
```

```
[208]: names.append('Income')
```

```
[209]: # Read the new file again
       df = pd.read_csv("adult_income_data.csv",names=names)
       df.head()
```

```
[209]:   age    workclass  fnlwgt  education  education-num  \
0    39    State-gov   77516    Bachelors              13
```


1	50	Self-emp-not-inc	83311	Bachelors	13
2	38	Private	215646	HS-grad	9
3	53	Private	234721	11th	7
4	28	Private	338409	Bachelors	13

	marital-status	occupation	relationship	sex \
0	Never-married	Adm-clerical	Not-in-family	Male
1	Married-civ-spouse	Exec-managerial	Husband	Male
2	Divorced	Handlers-cleaners	Not-in-family	Male
3	Married-civ-spouse	Handlers-cleaners	Husband	Male
4	Married-civ-spouse	Prof-specialty	Wife	Female

	capital-gain	capital-loss	hours-per-week	native-country	Income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

```
[210]: # Use the describe command to get the stastical summary of the dataset.
```

```
[211]: df.describe()
```

```
[211]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss \
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000

	hours-per-week
count	32561.000000
mean	40.437456
std	12.347429
min	1.000000
25%	40.000000
50%	40.000000
75%	45.000000
max	99.000000

```
[212]: # Make a list of all variable
```

```
[213]: # Make a list of all variables with classes
vars_class = ['workclass', 'education', 'marital-status',
```

```
'occupation','relationship','sex','native-country']
```

```
[214]: # Create a loop to count and print them by using the following command
```

```
[215]: for v in vars_class:
        classes=df[v].unique()
        num_classes = df[v].nunique()
        print("There are {} classes in the \"{}\" column. They are: {}".
        ↪format(num_classes,v,classes))
        print("-"*100)
```

```
There are 9 classes in the "workclass" column. They are: [' State-gov' ' Self-
emp-not-inc' ' Private' ' Federal-gov' ' Local-gov'
' ?' ' Self-emp-inc' ' Without-pay' ' Never-worked']
```

```
-----
There are 16 classes in the "education" column. They are: [' Bachelors' ' HS-
grad' ' 11th' ' Masters' ' 9th' ' Some-college'
' Assoc-acdm' ' Assoc-voc' ' 7th-8th' ' Doctorate' ' Prof-school'
' 5th-6th' ' 10th' ' 1st-4th' ' Preschool' ' 12th']
```

```
-----
There are 7 classes in the "marital-status" column. They are: [' Never-married'
' Married-civ-spouse' ' Divorced'
' Married-spouse-absent' ' Separated' ' Married-AF-spouse' ' Widowed']
```

```
-----
There are 15 classes in the "occupation" column. They are: [' Adm-clerical' '
Exec-managerial' ' Handlers-cleaners' ' Prof-specialty'
' Other-service' ' Sales' ' Craft-repair' ' Transport-moving'
' Farming-fishing' ' Machine-op-inspct' ' Tech-support' ' ?'
' Protective-serv' ' Armed-Forces' ' Priv-house-serv']
```

```
-----
There are 6 classes in the "relationship" column. They are: [' Not-in-family' '
Husband' ' Wife' ' Own-child' ' Unmarried'
' Other-relative']
```

```
-----
There are 2 classes in the "sex" column. They are: [' Male' ' Female']
```

```
-----
There are 42 classes in the "native-country" column. They are: [' United-States'
' Cuba' ' Jamaica' ' India' ' ?' ' Mexico' ' South'
' Puerto-Rico' ' Honduras' ' England' ' Canada' ' Germany' ' Iran'
' Philippines' ' Italy' ' Poland' ' Columbia' ' Cambodia' ' Thailand'
' Ecuador' ' Laos' ' Taiwan' ' Haiti' ' Portugal' ' Dominican-Republic'
' El-Salvador' ' France' ' Guatemala' ' China' ' Japan' ' Yugoslavia']
```

```
' Peru' ' Outlying-US(Guam-USVI-etc)' ' Scotland' ' Trinidad&Tobago'
' Greece' ' Nicaragua' ' Vietnam' ' Hong' ' Ireland' ' Hungary'
' Holand-Netherlands']
```

```
[216]: # Find the missing values by using the following command
```

```
[217]: df.isnull().sum()
```

```
[217]: age                0
workclass              0
fnlwgt                0
education              0
education-num          0
marital-status         0
occupation             0
relationship           0
sex                   0
capital-gain           0
capital-loss           0
hours-per-week         0
native-country         0
Income                0
dtype: int64
```

```
[218]: # Create a DataFrame with only age, education and occupation by using subnetting
```


```
[219]: df_subset = df[['age', 'education', 'occupation']]
df_subset.head()
```

```
[219]:   age  education  occupation
0   39  Bachelors  Adm-clerical
1   50  Bachelors  Exec-managerial
2   38   HS-grad  Handlers-cleaners
3   53    11th  Handlers-cleaners
4   28  Bachelors  Prof-specialty
```

```
[220]: # Plot a histogram of age with a bin size of 20
```


```
[221]: df_subset['age'].hist(bins=20)
```

```
[221]: <Axes: >
```



```
[222]: # PLOT box plots for age grouped by education
```

```
[223]: df_subset.boxplot(column='age',by='education',figsize=(25,10))
plt.xticks(fontsize=15)
plt.xlabel("Education",fontsize=20)
plt.show()
```



```
[224]: # create a function to strip the whitespace charecters
```

```
[225]: def strip_whitespace(s):  
        return s.strip()
```

```
[226]: # Use the apply method to apply the function to all the columns
```

```
[227]: import warnings  
        # Suppress all warnings  
        warnings.filterwarnings("ignore")  
  
        # Education column  
        df_subset['education_stripped'] = df['education'].apply(strip_whitespace)  
        df_subset['education'] = df_subset['education_stripped']  
        df_subset.drop(labels = ['education_stripped'],axis=1,inplace=True)  
  
        # Occupation column  
        df_subset['occupation_stripped'] = df['occupation'].apply(strip_whitespace)  
        df_subset['occupation'] =df_subset['occupation_stripped']  
        df_subset.drop(labels = ['occupation_stripped'],axis=1,inplace=True)
```

```
[228]: # Find the number of people who are aged between 30 & 50
```

```
[229]: # Conditional clauses and join them by & (AND)  
        df_filtered=df_subset[(df_subset['age']>=30) & (df_subset['age']<=50)]
```

```
[230]: # Check the contents of the new datasets
```

```
[231]: df_filtered.head()
```

```
[231]:   age  education  occupation  
0   39  Bachelors  Adm-clerical  
1   50  Bachelors  Exec-managerial  
2   38   HS-grad  Handlers-cleaners  
5   37   Masters  Exec-managerial  
6   49      9th    Other-service
```

```
[232]: # Find the shape of the filtered DataFrame and specify the index of the tuple_  
        ↪as 0 to return the first element
```

```
[233]: answer_1=df_filtered.shape[0]  
        answer_1
```

```
[233]: 16390
```

```
[234]: # Print the number of people of age between 30 and 50 in this dataset
```

```
[235]: print("There are {} people of age between 30 and 50 in this dataset.".
        ↪format(answer_1))
```

There are 16390 people of age between 30 and 50 in this dataset.

```
[236]: # Group by occupation and show the summary statistics by age
```

```
[237]: df_subset.groupby('occupation').describe()['age']
```

```
[237]:
```


	count	mean	std	min	25%	50%	75%	max
occupation								
?	1843.0	40.882800	20.336350	17.0	21.0	35.0	61.0	90.0
Adm-clerical	3770.0	36.964456	13.362998	17.0	26.0	35.0	46.0	90.0
Armed-Forces	9.0	30.222222	8.089774	23.0	24.0	29.0	34.0	46.0
Craft-repair	4099.0	39.031471	11.606436	17.0	30.0	38.0	47.0	90.0
Exec-managerial	4066.0	42.169208	11.974548	17.0	33.0	41.0	50.0	90.0
Farming-fishing	994.0	41.211268	15.070283	17.0	29.0	39.0	52.0	90.0
Handlers-cleaners	1370.0	32.165693	12.372635	17.0	23.0	29.0	39.0	90.0
Machine-op-inspct	2002.0	37.715285	12.068266	17.0	28.0	36.0	46.0	90.0
Other-service	3295.0	34.949621	14.521508	17.0	22.0	32.0	45.0	90.0
Priv-house-serv	149.0	41.724832	18.633688	17.0	24.0	40.0	57.0	81.0
Prof-specialty	4140.0	40.517633	12.016676	17.0	31.0	40.0	48.0	90.0
Protective-serv	649.0	38.953775	12.822062	17.0	29.0	36.0	47.0	90.0
Sales	3650.0	37.353973	14.186352	17.0	25.0	35.0	47.0	90.0
Tech-support	928.0	37.022629	11.316594	17.0	28.0	36.0	44.0	73.0
Transport-moving	1597.0	40.197871	12.450792	17.0	30.0	39.0	49.0	90.0

```
[238]: # Use subsets and groupby to find the outliers
```

```
[239]: occupation_stats= df_subset.groupby(
        'occupation').describe()['age']
```

```
[240]: # Plot the values on a bar chart
```

```
[241]: plt.figure(figsize=(15,8))
        plt.barh(y=occupation_stats.index,
                  width=occupation_stats['count'])
        plt.yticks(fontsize=13)
        plt.show()
```



```
[242]: #create new dataset where occupation is column. first create two such datasets
        ↳by taking random samples from the full datasets
```

```
[243]: df_1 = df[['age',
                'workclass',
                'occupation']].sample(5,random_state=101)
df_1.head()
```

```
[243]:      age workclass      occupation
22357   51   Private  Machine-op-inspct
26009   19   Private           Sales
20734   40   Private   Exec-managerial
17695   17   Private  Handlers-cleaners
27908   61   Private   Craft-repair
```

```
[244]: # The second dataset
```

```
[245]: df_2 = df[['education',
                'occupation']].sample(5,random_state=101)
df_2.head()
```

```
[245]:      education      occupation
22357   HS-grad  Machine-op-inspct
26009     11th           Sales
20734   HS-grad   Exec-managerial
17695    10th  Handlers-cleaners
27908   7th-8th   Craft-repair
```

```
[246]: # Merge the two datasets together
```

```
[247]: df_merged = pd.merge(df_1,df_2,
                             on='occupation',
                             how='inner').drop_duplicates()

df_merged
```

```
[247]:   age workclass      occupation education
0    51   Private  Machine-op-inspct   HS-grad
1    19   Private           Sales      11th
2    40   Private   Exec-managerial   HS-grad
3    17   Private  Handlers-cleaners    10th
4    61   Private   Craft-repair     7th-8th
```

```
[ ]:
```

```
[248]: # 3. Create a series and practice basic arithmetic steps
# a. Series 1 = 7.3, -2.5, 3.4, 1.5
# i. Index = 'a', 'c', 'd', 'e'
# b. Series 2 = -2.1, 3.6, -1.5, 4, 3.1
# i. Index = 'a', 'c', 'e', 'f', 'g'
# c. Add Series 1 and Series 2 together and print the results
# d. Subtract Series 1 from Series 2 and print the results
```

```
[249]: # Creating Series 1
series_1 = pd.Series([7.3, -2.5, 3.4, 1.5], index=['a', 'c', 'd', 'e'])
print("Series 1:")
print(series_1)
```

```
Series 1:
a      7.3
c     -2.5
d      3.4
e      1.5
dtype: float64
```

```
[250]: # Creating Series 2
series_2 = pd.Series([-2.1, 3.6, -1.5, 4, 3.1], index=['a', 'c', 'e', 'f', 'g'])
print("Series 2:")
print(series_2)
```

```
Series 2:
a     -2.1
c      3.6
e     -1.5
f      4.0
g      3.1
dtype: float64
```

```
[251]: # c. Add Series 1 and Series 2 together
```



```
[252]: # Adding Series 1 and Series 2
result_addition = series_1 + series_2
print("\nResult of adding Series 1 and Series 2:")
print(result_addition)
```

Result of adding Series 1 and Series 2:

```
a    5.2
c    1.1
d    NaN
e    0.0
f    NaN
g    NaN
dtype: float64
```

```
[253]: # d. Subtract Series 1 from Series 2
```

```
[254]: # Subtracting Series 1 from Series 2
result_subtraction = series_2 - series_1
print("\nResult of subtracting Series 1 from Series 2:")
print(result_subtraction)
```

Result of subtracting Series 1 from Series 2:

```
a   -9.4
c    6.1
d    NaN
e   -3.0
f    NaN
g    NaN
dtype: float64
```

```
[ ]:
```